

DB2锁问题处理最佳实践

徐明伟

北京普远天成科技有限公司 技术总监

DTCC2012

议题

1

DB2锁概述

2

DB2锁问题监控和定位

3

DB2锁问题调优

4

DB2 9.7锁机制深入分析

5

DB2锁案例分享

DTCC2012

为什么需要锁

- 一致性机制
 - 事务
 - 日志
 - 锁
 - 隔离级别
- 锁
 - 维护数据一致性
 - 控制并发性
- 锁分类
 - 锁的对象（表、行、表空间、索引）
 - 锁的模式（S,X等）

锁导致问题/锁现象

- 锁的几种现象
 - 锁等待
 - 锁超时
 - 死锁
 - 锁升级
 - 锁转换
- 锁产生的问题
 - 系统运行慢
 - 应用回滚

议题



1

DB2锁概述

2

DB2锁问题监控和定位

3

DB2锁问题调优

4

DB2 9.7锁机制深入分析

5

DB2锁案例分享

DTCC2012

锁问题监控和定位

- 锁问题监控定位工具
 - ✓ Snapshot快照
 - ✓ deadlock event monitor with details history
 - ✓ db2pd (8.2后)
 - ✓ db2pdcfg (9.1)
 - ✓ db2_capture_locktimeout (9.5)
 - ✓ New Locking event monitor (9.7)
- 锁是症状，不是根源

锁快照监控

- 通过get snapshot for database on <db> 或sysibmadm.snapdb
– 死锁、锁等、锁超时、锁升级等统计信息

Database Snapshot

Database name	= CRMDB
First database connect timestamp	= 10/10/2011 20:50:28.421301
Snapshot timestamp	= 03/07/2012 11:10:08.120847
Locks held currently	= 1
Lock waits	= 471967
Time database waited on locks (ms)	= 2039693347
Lock list memory in use (Bytes)	= 1740480
Deadlocks detected	= 470
Lock escalations	= 0
Exclusive lock escalations	= 0
Agents currently waiting on locks	= 0
Lock Timeouts	= 24782

数据库连接时间

快照时间

锁等数量

发生锁等的时间

死锁数量

锁升级数量

锁超时数量

DTCC2012

db2pd 监控锁等

- db2pd -d <db_name> -locks wait -tra -app -dyn
 - 定位引起锁等的事务、应用和动态语句

Locks:

Address	TranHdl	Lockname	Type	Mode	Sts	Owner
0x9A512080	11	030005000900000000000000052	Row	.NS	W	12
0x9A516400	12	030005000900000000000000052	Row	..X	G	12

Transactions:

Address	AppHandl	[mod-index] TranHdl	Locks	State
0x9A3E2F80	5899	[000-05899] 11	4	READ
0x9A3E3C80	5964	[000-05964] 12	3	WRITE

Applications:

Address	AppHandl	Status	C-AnchID	C-StmtUID	L-AnchID	L-StmtUID	Appid
0x20990060	5899	Lock-wait	321	1	85	1	*LOCAL.db2inst1.110126211438
0x20AD8430	5964	UOW-Waiting	0	0	734	8	*LOCAL.db2inst1.110126214851

Dynamic SQL Statements:

Address	AnchID	StmtUID	Text
0x9CB697D0	85	3	CALL REORGCHK IX_STATS ('T', 'DBA', '."SNAP_LOCKWAIT"')
0x9C9C9E20	321	1	←select * from t1 where coll='aaaa'
0x9CA09990	499	1	insert into t1 values('eeee','eeee')
0x7CDC0340	734	8	←select * from t1 where coll='eeee'

DTCC2012

db2pdcfg捕获锁超时 (9.1版本)

- db2pd结合db2cos回调脚本捕获锁超时或死锁
- 改写db2cos回调脚本

```
"LOCKTIMEOUT")
echo "Lock Timeout Caught"                >> $logfile
if [ ! -n "$database" ]
then
    db2pd -inst                            >> $logfile
else
    db2pd -db $database -locks -tra -app -dyn >> $logfile
fi
;;
```

- db2pdcfg -catch locktimeout count=1
- 当发生锁超时，会调用db2cos脚本

db2_capture_locktimeout捕获锁超时(9.1 fp4-9.5)

- 设置db2_capture_locktimeout注册变量
 - db2set DB2_CAPTURE_LOCKTIMEOUT=ON
- 创建死锁事件监控器
 - db2 “create event monitor dlockevm for deadlocks with details history write to file ‘/home/db2inst1/locks’ ”

```
Lock Requestor:          --锁请求者相关信息
  System Auth ID:        DB2INST1
  Application Handle:     [0-54]
  Application ID:         *LOCAL.db2inst1.110128054807
  Application Name:       db2bp
  Requesting Agent ID:    72
  Coordinator Agent ID:   72
  Coordinator Partition:  0
  Lock timeout Value:     15000 milliseconds
  Lock mode requested:    .NS
  Application Status:     (SQLM_UOWEXEC)
  Current Operation:      (SQLM_FETCH)
  Lock Escalation:        No

Context of Lock Request:
  Identification:         UOW ID (1); Activity ID (2)
  Activity Information:
    Package Schema:      (NULLID )
    Package Name:        (SQLC2H20NULLID )
    Package Version:     ()
    Section Entry Number: 201
    SQL Type:            Dynamic
    Statement Type:      DML, Select (blockable)
    Effective Isolation:  Cursor Stability
    Statement Unicode Flag: No
    Statement:            select * from t1 where coll='xxxx'
```

锁请求信息

请求锁模式

请求的SQL语句

DTCC2012

Lock Owner (Representative): --锁拥有者相关信息

System Auth ID: DB2INST1
Application Handle: [0-46]
Application ID: *LOCAL.db2inst1.110128054759
Application Name: db2bp
Requesting Agent ID: 19
Coordinator Agent ID: 19
Coordinator Partition: 0
Lock mode held: ..X

锁拥有者信息

--锁持有模式

当前锁模式

List of Active SQL Statements: Not available

List of Inactive SQL Statements from current UOW:

Entry: #1
Identification: UOW ID (3); Activity ID (2)
Package Schema: (NULLID)
Package Name: (SQLC2H20)
Package Version: ()
Section Entry Number: 201
SQL Type: Dynamic
Statement Type: DML, Select (blockable)
Effective Isolation: Cursor Stability
Statement Unicode Flag: No
Statement: select * from t1 where coll= ffff' --SQL 语句

Entry: #2
Identification: UOW ID (3); Activity ID (1)
Package Schema: (NULLID)
Package Name: (SQLC2H20)
Package Version: ()
Section Entry Number: 203
SQL Type: Dynamic
Statement Type: DML, Insert/Update/Delete
Effective Isolation: Cursor Stability
Statement Unicode Flag: No
Statement: insert into t1 values('ffff','ffff') --SQL 语句

占有锁的SQL语句

DTCC2012

Deadlock event monitor 监控死锁

• 创建deadlock事件监控器

- db2 "create event monitor dlockevm for deadlocks with details history write to file '/home/db2inst1/deadlock' "

```
5) Deadlocked Connection ...
Deadlock ID: 1
Participant no.: 2
Participant no. holding the lock: 1
Appl Id: *LOCAL.db2inst1.110128063201
Appl Id of connection holding the lock: *LOCAL.db2inst1.110128054759
Seq. no. of connection holding the lock: 00001
Lock wait start time: 01/27/2011 22:37:04.601677
Lock Name : 0x030005000400000000000000000052
Current Mode : none
Deadlock detection time: 01/27/2011 22:37:09.788844
Table of lock waited on : T1
Schema of lock waited on : DB2INST1
Type of lock: Row
Mode of lock: X - Exclusive
Mode application requested on lock: NS - Share (CS/RS)
Deadlocked Statement:
Type : Dynamic
Operation: Fetch
Text : select * from t1

6) Deadlock statement history ...
Deadlock ID : 1
Participant No : 2
Stmt history ID : 2
Type : Dynamic
Statement text : select * from t1

7) Deadlock statement history ...
Deadlock ID : 1
Participant No : 2
Stmt history ID : 1
Type : Dynamic
Statement text : insert into t2 values('bbb')
```

Deadlock id

死锁的表

表上已有的锁模式

当前请求的锁模式

死锁语句

DTCC2012

Deadlock event monitor 监控死锁(2)

9) Deadlocked Connection ...

Deadlock ID: 1
Participant no.: 1
Participant no. holding the lock: 2
Appl Id: *LOCAL.db2inst1.110128054759
Appl Seq number: 00018
Appl Id of connection holding the lock: *LOCAL.db2inst1.110128063201
Seq. no. of connection holding the lock: 00001
Lock wait start time: 01/27/2011 22:37:03.515204
Lock Name : 0x030006000500000000000000000052
Current Mode : none
Deadlock detection time: 01/27/2011 22:37:09.795276
Table of lock waited on : T2
Schema of lock waited on : DB2INST1
Type of lock: Row
Mode of lock: X - Exclusive
Mode application Requested on lock: NS - Share (CS/RS)
Deadlocked Statement:
Type : Dynamic
Operation: Fetch
Text : select * from t2

Deadlock id

锁的表

请求的锁类型

10) Deadlock statement history ...

Deadlock ID : 1
Participant No : 1
Stmt history ID : 2
Type : Dynamic
Statement text : select * from t2

11) Deadlock statement history ...

Deadlock ID : 1
Participant No : 1
Stmt history ID : 1
Type : Dynamic
Statement text : insert into t1 values('aaa')

死锁语句

DTCC2012

锁事件监控器(9.7版本)

- 9.7引入了新的锁事件监控器
- 用同一个锁事件监控器就可捕获死锁、锁超时和锁等待
- 采用UE表存取锁事件结果，使得分析更简单
- 锁事件监控器的使用包含三个步骤：
 - 1. 创建锁事件监控器
 - 2. 设置数据收集的类型和级别
 - 3. 格式化并分析数据

创建锁事件监控器

create event monitor lockevmon **for locking**
write to unformatted event table (table
locks)

创建锁事件
监控器

创建锁事件
监控器

set event monitor lockevmon state=1

- Unformatted Event(UE)表
 - 解释说明

DTCC2012

设置锁事件参数

- 捕获死锁事件
 - `db2 update db cfg using mon_deadlock hist_and_values`
- 捕获锁超时事件
 - `db2 update db cfg using mon_locktimeout hist_and_values`
- 捕获锁等待事件
 - `db2 update db cfg using mon_lw_thresh 5000000`
 - `db2 update db cfg using mon_lockwait hist_and_values`
- 要监控的锁事件信息详细程度:
 - `Without_hist` : 收集基本事件信息
 - `With_hist`: 一个事务里最多收集250个活动
 - `Hist_and_values`: 收集活动和值
 - `None`: 不收集事件

格式化锁事件输出结果

- 一旦捕获了锁事件，下一步就要对锁数据进行分析
- 对UE表数据的格式化有三种方法：
 - Db2evmonfmttool
 - ✓ 产生文本格式输出报告
 - EVMON_FORMAT_UE_TO_XML_UDF函数
 - ✓ 产生XML格式输出报告
 - EVMON_FORMAT_UE_TO_TABLE存储过程
 - ✓ 将UE数据格式化成关系表数据
 - ✓ 分析处理变得简单

死锁举例

Application #1

```
update t2 set name='t2...' where id=200  
select * from employee
```

Application #2


```
update t1 set name='t11...' where id=100  
update t2 set name='t2new...' where id=200
```

```
update t1 set name='t222' where id=100 (911 error)
```

模拟一个死锁场景

使用EVMON_FORMAT_UE_TO_TABLE格式化

```
call EVMON_FORMAT_UE_TO_TABLES(  
'LOCKING',NULL,NULL,NULL,NULL,NULL,'RECREATE_FO  
RCE',-1, 'select * from locks order by event_timestamp'  
)
```



LOCK_ACTIVITY_VALUES	INST97	T	2012-04-03-22.22.42.368824
LOCK_EVENT	INST97	T	2012-04-03-22.22.41.600176
LOCK_PARTICIPANTS	INST97	T	2012-04-03-22.22.41.921119
LOCK_PARTICIPANT_ACTIVITIES	INST97	T	2012-04-03-22.22.42.112358

查看死锁和参与者信息

```
Select xmlid, event_id, event_type,  
event_timestamp, member, dl_conns,  
rolled_back_participant_no from lock_event
```

查看锁信息

XMLID	EVENT_TYPE	EVENT_TIMESTAMP	MEMBER	DL_CONNS	ROLLED_BACK_PARTICIPANT_NO
db2LockEvent_4...	DEADLOCK	2012-04-03-21.27.19.660632	0	2	2

```
Select xmlid, participant_no, participant_type,  
participant_no_holding_lk, application_handle from  
lock_participants where xmlid like db2LockEvent_1%'
```

查看死锁参
与者信息

XMLID	PARTICIPANT_NO	PARTICIPANT_TYPE	PARTICIPANT_NO_HOLDING_LK	APPLICATION_HANDLE
db2LockEvent_4...	1	Requester	2	7
db2LockEvent_4...	2	Requester	1	60

DTCC2012

查看死锁参与者SQL活动信息

```
select activity_id, activity_type, uow_id, stmt_type,  
       substr(stmt_text,1,100)as stmt_text  
from lock_participant_activities  
where xmlid like 'db2LockEvent_4%'
```

查看死锁参与者SQL活动信息

ACTIVITY_ID	ACTIVITY_TYPE	UOW_ID	STMT_TYPE	STMT_TEXT
1	past	4	2	update t1 set name='t11...' where id=100
2	current	4	2	update t2 set name='t2new...' where id=200
1	past	1	2	update t2 set name='t2...' where id=200
2	past	1	2	select * from employee
3	current	1	2	update t1 set name='t222' where id=100

议题

1

DB2锁概述

2

DB2锁问题监控和定位

3

DB2锁问题调优

4

DB2 9.7锁机制深入分析

5

DB2锁案例分享

DTCC2012

性能调优关键

- 锁是症状，SQL是问题的根源
- I/O 最关键
 - 减少I/O
 - 最大化I/O效率
 - 存储规划
 - 物理设计
- CPU 2大杀手
 - 表扫描
 - 排序
- Memory命中率可能会骗人

锁的调优(1) — 应用层

- 写优秀的SQL语句
- 创建合适的索引避免表扫描
- 选择合适的隔离级别: UR,CS(CC),RS,RR
- 事务尽可能频繁的提交
- 在事务结尾执行insert/update/delete

锁的调优(2) —数据库层

- 调优locklist 和maxlocks数据库参数
 - 9.1后设为automatic
- Locktimeout 锁超时参数
 - OLTP系统, 建议设置为15-30秒
 - DW系统,建议60-120秒
- 锁延迟设计考虑
 - db2set DB2_EVALUNCOMMITTED=ON
 - db2set DB2_SKIPDELETED=ON
 - db2set DB2_SKIPINSERTED=ON
- 9.7 当前已提交读特性(currently committed)
- 数据表维护
 - Runstats更新统计信息
 - Reorg重组表和索引碎片
 - Rebind重新绑定包, 更新执行计划

写好的SQL语句

- 只返回需要的行，避免用 `select * from t1`
- 加过滤条件限制返回的行数
- 避免笛卡尔乘积， `select * from a,b`
- 使用参数化查询， `where col1=?`，减少编译时间
 - 案例分析
- 避免对查询条件计算， `where salary*2>xx` 改为 `salary > xx/2`
 - 无法利用索引
- 使用 `for read only` 或 `for fetch only`
- 使用 `for update of`
- 避免数字类型转换
- 避免数据类型不匹配
- 如果可能，尽量避免使用 `order by` 和 `distinct`
- 尽量使用 `exists` 而不是用 `in`
- 函数的效率很高，充分利用

DTCC2012

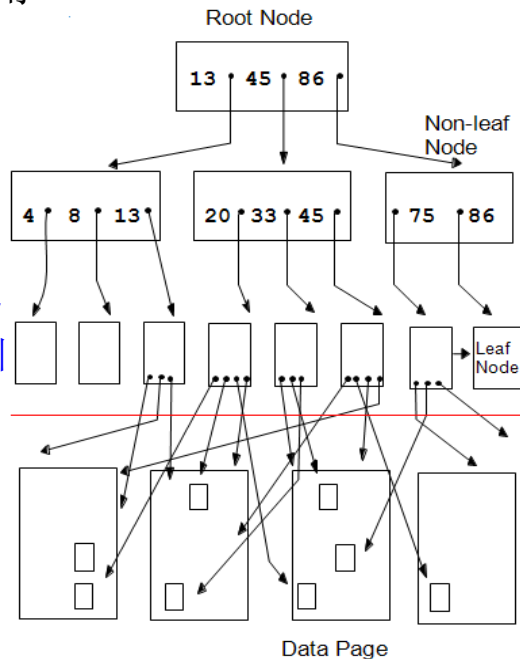
创建索引

- 索引对于提高SQL读有无可替代作用

- 提高查询性能
- 避免不必要的表扫描
- 减少排序、减少锁
- 减少CPU和I/O使用

- 索引创建最佳实践

- 为where查询条件、Sort排序(order by、max()、min ()等)、join谓词创建索引
- 分析组合索引键的顺序
 - (a,b), (b,a) 完全不同
- 不要创建冗余索引
 - (a), (a,b) (a)为冗余索引
- 确保索引被用到
- 创建cluster 索引
- 通过include语句创建index-only索引
- 对于不稳定数据, 可考虑用volatile强制走索引
- 推荐使用db2advis建议索引



物理设计

- 海量数据库物理设计
 - 索引
 - 多维索引
 - 物化视图
 - 分区表
 - 压缩
 - 数据库分区
- 数据归档
 - 减少数据，减少I/O等访问

议题

1

DB2锁概述

2

DB2锁问题监控和定位

3

DB2锁问题调优

4

DB2 9.7锁机制深入分析

5

DB2锁案例分享

DTCC2012

隔离级别(isolation level)

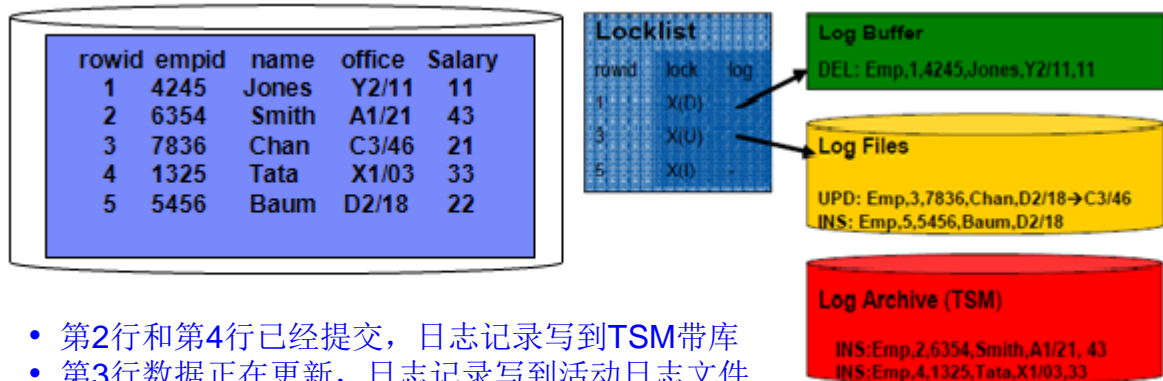
- 隔离级别控制锁的范围和粒度，**隔离级别只适用于读**

DB2 Isolation	ANSI Isolation	Dirty Write	Dirty Read	Fuzzy Read	Phantom Read
Uncommitted Read (UR)	Read Uncommitted (Level 0)	×	✓	✓	✓
Cursor Stability (CS)	Read Committed (Level 1)	×	×	✓	✓
Read Stability (RS)	Repeatable Read (Level 2)	×	×	×	✓
Repeatable Read (RR)	Serializable (Level 3)	×	×	×	×

- DB2 9.7版本引入了currently committed机制
 - ✓ 如果发现未提交的行，则使用当前已经提交的数据(before image)
 - ✓ 写不阻碍读，减少锁等和死锁，超时等
 - ✓ 在locklist锁信息中加入标识，当另外事务读取数据时基于标识判断
 - ✓ 通过cur_commit数据库参数控制

DTCC2012

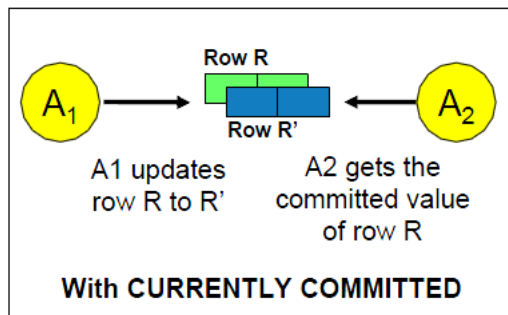
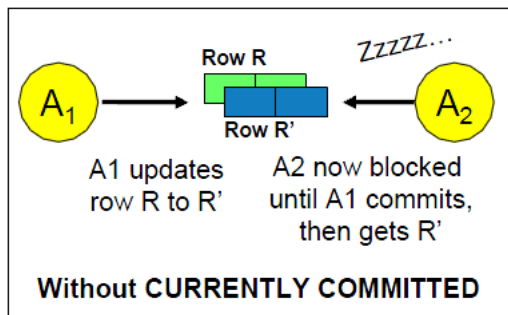
Currently Committed机制



- 第2行和第4行已经提交，日志记录写到TSM带库
 - 第3行数据正在更新，日志记录写到活动日志文件
 - 第5行数据正在插入，日志记录写到活动日志文件
 - 第1行数据正在删除，日志记录在Log buffer内存
-
- 当启用了CC机制，DB2会在每个行锁上增加一个标识：
 - No information: 表示记录已经加锁。
 - Uncommitted insert identifier: 表示这行是新插入未提交行。
 - Log information: 表示这行没有提交，包含了当前已落实数据的日志LSN。

DTCC2012

Currently Committed性能影响



- 使用currently committed的tips
 - 如果观察到大量读写锁等待时间，使用cc更好
 - 严重读写竞争的应用会有5-10%以上的性能提升
 - 建议调整logbufsz为2048，9.7缺省是256页

议题

1

DB2锁概述

2

DB2锁问题监控和定位

3

DB2锁问题调优

4

DB2 9.7锁机制深入分析

5

DB2锁案例分享

DTCC2012

案例1：某通信行业电子运维系统

• 问题概述：

- 全省统一的电子运维管理系统
- 1000多个并发，高峰期业务无法处理
- 大量锁等，响应慢，CPU使用率80%以上

• 数据库监控：

First database connect timestamp	= 10/10/2011 20:50:28.421301
Snapshot timestamp	= 11/09/2011 17:19:31.409497

Number of Threshold Violations	= 0
Locks held currently	= 4
Lock waits	= 86248
Time database waited on locks (ms)	= 1306986473
Lock list memory in use (Bytes)	= 2186752
Deadlocks detected	= 222
Lock escalations	= 0
Exclusive lock escalations	= 0
Agents currently waiting on locks	= 0
Lock Timeouts	= 15560

First database connect timestamp	= 11/10/2011 20:09:23.982
Snapshot timestamp	= 12/11/2011 18:32:05.902

Number of Threshold Violations	= 0
Locks held currently	= 4
Lock waits	= 9429
Time database waited on locks (ms)	= 183176718
Lock list memory in use (Bytes)	= 198742
Deadlocks detected	= 42
Lock escalations	= 0
Exclusive lock escalations	= 0
Agents currently waiting on locks	= 0
Lock Timeouts	= 2305

• 解决办法：

- 通过db2pd找到占有锁的SQL语句，优化
- 调优后，死锁和锁超时数据降低了80%，性能大大提升

DTCC2012

案例2：北京某部委优化项目

• 问题概述：

- 北京市人口统计信息统一管理系统
- 平常800个并发，高峰1000多个，高峰期查询需要10分钟
- 大量锁等，响应慢，CPU利用率低

• 数据库监控：

First database connect timestamp	= 01/16/2012 19:23:39.036989
Snapshot timestamp	= 02/17/2012 13:45:47.147392
Locks held currently	= 263
Lock waits	= 2772575
Time database waited on locks (ms)	= 9986324507
Lock list memory in use (Bytes)	= 49920
Deadlocks detected	= 4
Lock escalations	= 0
Exclusive lock escalations	= 0
Agents currently waiting on locks	= 1
Lock Timeouts	= 0

说明：

平均每天锁等达到 **90万个**
平均每个锁等等等待 **3.6秒**
累计锁等时间达到 **115天**
锁超时次数为**0**?

• 解决办法：

- 调优bufferpool，从1.5G到25G
- 调优了5条SQL语句
- 调优后，性能大幅度提升，最慢的响应时间从10分钟降为秒级

DTCC2012

案例3：某银行核心交易死锁问题处理(1)

- 问题概述：

- 核心银行交易系统
- 12月25日调优后，性能下降严重，死锁数量急剧增加

- 数据库监控：

First database connect timestamp	= 12/29/2011 00:13:12.846036
Snapshot timestamp	= 12/30/2011 11:50:17.975290
Locks held currently	= 263
Lock waits	= 62245
Time database waited on locks (ms)	= 10514064
Lock list memory in use (Bytes)	= 254720
Deadlocks detected	= 531
Lock escalations	= 0
Exclusive lock escalations	= 0
Agents currently waiting on locks	= 0
Lock Timeouts	= 10

说明：

1天半时间发生了**531**个死锁

- 分析：

- 应用是通过sqlc开发，嵌入式静态语句
- 创建deadlock event monitor，发现有一条语句存在重大嫌疑
- 该语句访问的表数据量大概20万行，单条语句执行速度很快，大概2秒

案例3：某银行核心交易死锁问题处理(2)

- 通过db2expln查看执行计划，发现该语句cost高达9645

```
Statement:
DECLARE BIGAMT_CUR CURSOR
FOR
  SELECT *
  FROM dpsbigamtrans
  WHERE trandate =:H01057 AND acctno =:H01069 AND subacct
=:H01071
  WITH UR
```

Statement Isolation Level = Uncommitted Read

Section Code Page = 1386

Estimated Cost = 9645.297852

Estimated Cardinality = 0.000002

```
Access Table Name = VBSRUN.DPSBIGAMTTRANS ID
= 3,3337
| #Columns = 26
| Relation Scan
| | Prefetch: Eligible
| Lock Intents
| | Table: Intent Share
| | Row : Next Key Share
| Sargable Predicate(s)
| | #Predicates = 3
| | Return Data to Application
| | | #Columns = 26
Return Data Completion
```

End of section

Optimizer Plan:

```
RETURN
( 1)
|
TBSCAN
( 2)
|
Table:
VBSRUN
DPSBIGAMTTRANS
```

- 解决办法：
 - 创建索引
 - Runstats
 - Rebind package
 - 锁数量大大降低

DTCC2012

联系方式



- 徐明伟

- 资深DB2顾问（咨询、培训、服务支持）
- 手机：13701141650
- Email: xumwdb2@163.com
- qq: 907342263
- 新浪微博：徐明伟 db2咨询

- 取得成绩：

- 2012年“IBM DB2迁移之星大赛”冠军团队
- 2011年，出版《DB2数据库管理最佳实践》书籍
- 2011年“IBM DB2十大人物”
- 2010年“IBM软件技术精英年度成就会员”

DTCC2012



Thank
You

DTCC2012